

Task 2: Stream Processing with Spark Streaming

Platform setup, Dataset analysis, Cleaning the dataset

To perform Task 2, I set up a real-time data pipeline with Kafka v0.8.2, Spark Streaming v1.3.1 and Cassandra v2.1. Scala platform is v2.10 and Java v1.7. For further details about my clustered platform setup and cleaning the dataset, please refer to my [Task 1 report](#).

Exploring Spark and Spark Streaming

Approaching the new Task, I installed and configured Zeppelin tool in Ambari and played with its UI to explore Spark and Spark Streaming API.

Zeppelin integrates all in one Spark console and many other tools, among all, Kafka and Cassandra. I managed to automate some tasks, like loading libraries in the classpath and preparing HDFS filesystem and Cassandra, and run my first Spark scala commands, test the behaviour of some transformations and, more important, testing how Spark Cassandra connector works and effectively stores data in Cassandra!

Java Spark

After many attempts, a few successful, I decided that I needed to write more advanced code to make the best possible from Spark streaming transformations and I moved to Spark Java.

Kafka

I used Kafka to stream the data source. I set up one Kafka broker on purpose, local to one of my cluster nodes, and created one topic, *aviation_ontime*.

I used Kafka-console-producer shell script, with awk converting from a simple message to a keyed-message.

```
hadoop fs -cat /cloudcapstone/sourcedata/ontime_CLEAN/*/part-* | awk -F "\t" '{print $5,"$0"} | kafka-console-producer.sh --broker-list node2:6667 --property parse.key=true --property key.separator=, --topic aviation_ontime && echo -e "\n"
```

Cassandra

[Spark Cassandra connector](#) is my choice to process the data stream and send the results to the store. I successfully learned how to configure and use it to update the store with partial results, each streaming round, in a real-time streaming fashion.

Parallelism

Most of my time for Task 2 was spent in exploring Spark API. With all tools in a pipeline, it is very important to set up the parallelism level and be able to scale out the whole process.

Kafka is very flexible in creating several partitions within a topic. The partitions are necessary to allow a consumer to use more threads in reading from a stream. I applied the [directStream](#) approach, where Spark creates simply as many RDD partitions as Kafka partitions and reads data in parallel. Then, it's just a matter of defining the right key for your RDD transformation and all your data will be processed in parallel.

Dataset preparation for a real-time streaming

After the Task 1 I have in my HDFS filesystem the source data cleaned, which I can use for Task 2.

To solve the problem G3Q2 in the best possible way, I prepared a specific dataset for year 2008 in which I had a stream of events (flights) ordered by date. This is not a constraint to my process, but it makes same optimizations more effective, as I discuss them later on in this document.

Perhaps you would agree that such an assumption of receiving an ordered data stream is appropriate for a streaming process. According to my concept of data streaming pipeline, I strongly wanted to work in a real-time streaming fashion, as much as possible. This is the reason why I stressed my solution to compute and store partial results in real-time.

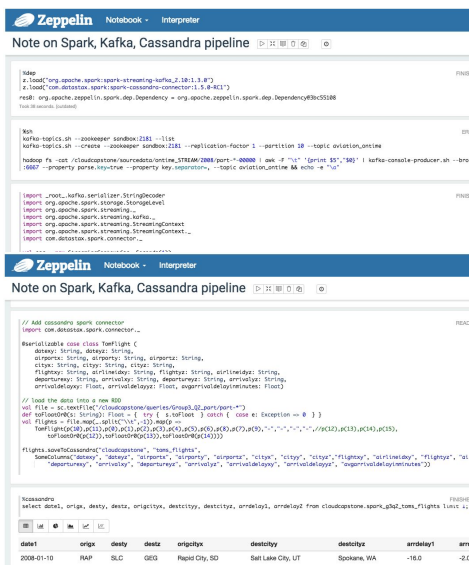
Solving the queries with Spark Streaming

I developed 2 Spark jobs:

G12StreamingJob. It groups together 5 processors for the 5 queries of groups 1&2 (G1Q1, G1Q3, G2Q1, G2Q2, G2Q4). This job runs the 5 processors in parallel. Each processor is attached to a dedicated stream and receives data from each of Kafka's topics partition.

G3StreamingJob Usage: <kafka-broker-host:port> <kafka-topic> <output-path> <cassandra-host> [<stream-freq-in-secs(def:1)> <cassandra-concurrent-writes> <cassandra-batch-size>]

G3StreamingJob. A "stateful" job, optimized and dedicated to the G3Q2 problem.



```
/usr/hdp/current/kafka-broker/bin/kafka-topics.sh --zookeeper sandbox:2181 --describe --topic aviati

Topic:aviation_ontime PartitionCount:10 ReplicationFactor:1 Configs:retention.ms=3600000
Topic:aviation_ontime Partition: 0 Leader: 0 Replicas: 0 Isr: 0
Topic:aviation_ontime Partition: 1 Leader: 0 Replicas: 0 Isr: 0
Topic:aviation_ontime Partition: 2 Leader: 0 Replicas: 0 Isr: 0
Topic:aviation_ontime Partition: 3 Leader: 0 Replicas: 0 Isr: 0
```

| date1 | origo | desty | destr | origcityp | destcityp | destcity | arrdelay1 | arr |
|------------|-------|-------|-------|----------------|--------------------|-------------|-----------|-----|
| 2008-01-10 | IRAP | SLC | OEK | Rapid City, SD | Salt Lake City, UT | Spokane, WA | -18.0 | -24 |

G3StreamingJob Usage: <kafka-broker-host:port> <kafka-topic> <output-path> <cassandra-host> [<stream-freq-in-secs(def:1)> <purge-days(def:31)> run-for-days(def:365) <cassandra-concurrent-writes> <cassandra-batch-size>]

I let some parameters to make the jobs configurable. Among all, please notice the frequency of streaming, which sets the batch duration in Spark Streaming context, and some optional settings to improve the performance of Spark Cassandra connector.

In addition, I developed a *KafkaConsumer* to help with the creation of the stream. The current implementation uses a direct stream ([KafkaUtils.createDirectStream](#)). I tried also the approach with receivers, but I noticed no difference and opted for the one described in Spark documentation as a more simple and efficient solution.

The main difference between my two jobs is how they accumulate and store the results.

Accumulators vs. UpdateStateByKey

Both jobs process a bit of the stream for each period and compute the partial results. But while the G12 job uses Spark concept of [Accumulator](#), collecting them all and flush them to disk and to datastore at frequent intervals (which I called *savepoints*), the G3 job uses Spark [updateStateByKey](#) operation.

This operation allows to transform the stream into a stateful RDD that can be updated by a user defined function. The RDD is of course partitioned and distributed, like Accumulator is, but the latter has got the drawback of collecting all the “local” values into a global one (centralized) before storing them. *UpdateStateByKey* implicitly uses *Checkpointing* to be fail safe.

With a stateful RDD, I was able to combine in memory all the flights and, finally, store the partial results into the datastore as soon as the streaming process was producing them in real-time! I designed some optimizations to keep in memory only the relevant portion of flights. For more details, please see below the description of the G3Q2 problem.

Please refer to my [Task 1 report](#) for the definitions of: *On-time arrival/departure performance*, *Average arrival/departure delay*, *Diverted/Cancelled flights*.

Group 1 Q1

G12StreamingJob uses *G1Q1Processor* to process the data and accumulate the results into *G1Q1Accumulable*.

This processor uses the following transformations to get the result:

- parses the incoming flights’ data, including the cancelled/diverted ones, and for each flight, it applies a mapping to create a pair RDD with *key = origin airport* and a second pair RDD with *key = destination airport*;
- makes a union of the two RDDs and reduces it by key, which gives the count of IN/OUT flight from every airport;

The final RDD *countOfFlightsByAirport* is still partitioned by key, locally to all the executors. Each partition can be processed and the data added to the shared accumulator.

The count of flights is increased in each round.

At the next savepoint, the accumulator value is collected from the master process, and a partial result is ready to be stored in HDFS and Cassandra. In my run, a savepoint usually happens every 30 secs.

Group 1 Q3

G12StreamingJob uses *G1Q3Processor* and produces results into the shared *G1Q3Accumulable*. See the screenshot of my result (*day, on-time performance, #flights, mean delay*).

This processor uses the following transformations:

- parses the incoming flights’ data, skipping the cancelled/diverted ones, and for each flight, it applies a mapping to create a pair RDD by *key = day of the week* and value including the arrival delay, if positive;
- reducing the pair RDD by key, it is possible to count the number of flights, count the flights on-time and compute the on-time arrival performance, finally, summing the arrival delays and computing a partial average delay.

The final RDD *perfByWeekDay* is processed for each partition and the result added to the accumulator. The count of flights and the arrival delay are increased, the mean arrival delay and the on-time arrival performance computed in each round.

A savepoint occurs after a configured number of periods, then the partial results are all saved to HDFS and stored in Cassandra.

Group 2 Q1

G12StreamingJob uses *G2Q1Processor* and *G2Q1Accumulable* to process the stream of flights, using the following transformations:

- it parses the flights, skipping the cancelled/diverted ones, and mapping a pair RDD with *key = origin airport + "/" + airline* and value including the departure delay, if positive;

| | |
|--|-----------------|
| / cloudcapstone / spark-results / G1Q1 | |
| G1Q1 | |
| Name | (ORD, 12330165) |
| .. | (ATL, 11399905) |
| | (DFW, 10703206) |
| | (LAX, 7651102) |
| | (PHX, 6520744) |
| | (DEN, 6191040) |
| | (DTW, 5584103) |
| | (IAH, 5418452) |
| | (MSP, 5157201) |
| | (SFO, 5122411) |
| | (STL, 5103717) |
| | (EWR, 5091532) |
| | (LAS, 4905679) |
| | (CLT, 4781947) |
| | (LGA, 4296328) |
| | (BOS, 4271528) |
| | (PHI, 4045177) |
| _SUCCESS | |
| Updated 2016-02-17 00:54 | |
| part-00000 | |
| Updated 2016-02-17 00:54 | |

| | |
|--------------------------------------|--|
| cloudcapstone / spark-results / G1Q3 | |
| G1Q3 | |
| Name | (Friday, 0.4720069, 16563498, 13.568059) |
| . | (Thursday, 0.48142144, 16525721, 13.0030775) |
| | (Wednesday, 0.5150601, 16531138, 11.463943) |
| | (Monday, 0.52942944, 16619491, 11.096879) |
| | (Sunday, 0.53240126, 15755185, 11.23013) |
| | (Tuesday, 0.5374497, 16513736, 10.492088) |
| | (Saturday, 0.57176596, 14614235, 9.4440975) |
| part-00000 | |
| Updated 2016-02-17 00:54 | |
| part-00000 | |
| Updated 2016-02-17 00:54 | |

- The *performanceRank* RDD gives for every airport, 10 carriers in descending order of on-time departure performance. Each partition is processed and aggregated into the accumable object, which will be stored in HDFS and Cassandra at the next savepoint.

| airport | ontimedepartureperf | airline | avgdeparturedelayinminutes | carrier | city | countofflights |
|---------|---------------------|---------|----------------------------|---------|-------------|----------------|
| SEA | 0.832117 | 20378 | 7.30657 | YV | Seattle, WA | 548 |
| SEA | 0.748656 | 20304 | 4.70382 | 00 | Seattle, WA | 27341 |
| SEA | 0.689221 | 20312 | 8.38985 | TZ | Seattle, WA | 2542 |
| SEA | 0.680443 | 19707 | 11.46716 | EA | Seattle, WA | 4065 |
| SEA | 0.666212 | 19386 | 7.21636 | NW | Seattle, WA | 118000 |
| SEA | 0.651852 | 20366 | 8.37037 | EV | Seattle, WA | 270 |
| SEA | 0.651282 | 20404 | 12.14359 | DH | Seattle, WA | 195 |
| SEA | 0.647513 | 19690 | 8.16441 | HA | Seattle, WA | 3759 |
| SEA | 0.630984 | 19805 | 7.32507 | AA | Seattle, WA | 123000 |
| SEA | 0.628542 | 19991 | 9.17828 | HP | Seattle, WA | 46051 |

G12StreamingJob uses *G2Q2Processor* and *G2Q2Accumulable* to process the stream of flights. This process is very similar to the one described in *G2Q1Processor*, with the only difference in the key used to map and reduce the pair RDD.

G12StreamingJob uses *G2Q4Processor* and *G2Q4Accumulable* to process the data and accumulate the results, by using the following transformations:

- | origin | destination | avgarrivaldelayinminutes | countofflights |
|--------|-------------|--------------------------|----------------|
| LGA | BOS | 7.66222 | 166438 |

| origin | destination | avgarrivaldelayinminutes | countofflights |
|--------|-------------|--------------------------|----------------|
| BOS | LGA | 8.49688 | 163647 |

G3StreamingJob uses the *G3Q2Processor* to process the data stream. There is no accumulator. Instead, a state RDD is computed every round and produces new combinations and updates, soon saved in the store.

- the population of incoming flights, excluding the cancelled/diverted ones, is split into two RDD *flightsLeg1*, *flightsLeg2*;
- the two RDD contains multiple solutions for the same day and route (X-Y), then we can reduce by (*key* = *date* + "|" + *orig* + "|" + *dest*) and keep only the best flights, e.g. the ones with the lower arrival delay;
- at this point, the best flights from both legs can be mapped to a tuple with (*key* = *date1* + "|" + *destY*) and (*value* = *FlighCombi*) - *FlighCombi* model is a composition of [*<leg1>*, *<leg2>*] matching the constraints of the same intermediate airport Y and the 2-day stop for sightseeing;

From that point, a state RDD is computed every round, by using the operation *updateStateByKey* with the function *mergeState*. The tuples in the state RDD, which are processed by the *mergeState* function are of type *G3Q2State*, which contains a *Map<String, Set<FlightCombi>>*, where the key is always *(date1 + "/" + destY)*. The given function processes the partitions of the state RDD to create new combinations and update the current combinations with better legs, eventually.

The state RDD grows in memory but it remains well distributed locally to the executors. The combinations computed in the state are sent to the data store at the end of every batch period. This real-time flowing of all the updates could result into a dangerous overloading of Cassandra. At this point I needed to design some optimizations to simplify the process and avoid the stress of the resources.

First of all, I chose `key = date1 + "/" + destY`, which allows the best partitioning of the flights. With the strict requirement of combining flights where the leg-2 is exactly 2 day after the leg-1, I was able to transform the `date2` of a leg-2 flight into the `date1` required by the key. This way, all the matching leg-1 and leg-2 are in the same set of *FlightCombi*, partitioned close together! Then, I introduced some useful optimizations, here described.

Flight events ordered, Time-window and Purge

If we assume a **partial/total ordering** of the events (thanks for introducing me to the concept, prof. Indy :-), it would be possible for Spark job to keep for each partition only the last two days of the streamed data, because for that partition (*key = date + "/" + destY*), there will be no incoming event out the date range [*date, date+2d*].

But total ordering is hard to achieve with Kafka broker [using many partitions](#), we can instead try to achieve a partial order in an extended **time-window** (i.e. 10 days or a month).

For this reason I made the assumption that the stream should carry all the flights with the same date as much as possible **close together**, or at least **within a time window**. To put it simple, defining in my process a window of 31 days, it would mean that while the flights with date 01/04/2008 are streamed, it is still allowed that a flight with date 01/03/2008 comes late to be processed. But any flight before that date would not be processed correctly!

Then, to put it in practice, I simply ordered the events by date, which is a very plausible assumption for a real-time streaming system.

I designed algorithm to compute, at every round, the date of the **less recent** incoming event (*dateLB*) and define a configurable number of day, i.e 31. Then I can calculate *purgeDate* = (*dateLB* - 31) to purge the old combinations older than *purgeDate* (no more useful and already stored in Cassandra). That is! At the end of the round, I simply discard from the memory all the partial combinations out of the time-window [*purgeDate, ∞*]. In the next round, the state will be smaller, but still able to combine the incoming flights, because I assumed that they will have a date not before the purged window lower bound!

As you can see, at this point it's just a matter of ordering the events, configure the right time-window, together with using a limited number of Kafka partition, which should not shuffle the data too much... and we have got a smarter process producing results 100% correct!

Storing updates only

To reduce the load in the store, I optimized the job to flag a *FlightCombi* when it is created or updated, and to store into Cassandra only the flagged ones (check the video to see how low is the resulting write requests rate).

Kafka optimization

I produced keyed messages to make the broker maintain the order of the stream, at least within each partition. With *directStream* approach, there is no need to use groups for the consumers, they all will receive the same messages queued in the one topic.

Cassandra optimization

Storing with savepoints vs. storing in real-time really makes the difference, with the first solution to avoid overloading the store with a lot of partial updates. You can see in this screenshot a early run of G3StreamingJob. Please notice the huge commit log of one of my Cassandra nodes, with my initial unoptimized G3Q2

implementation. At the beginning, I was favouring the first option, until I [finally optimized the G3Q2 solution](#) to store only the updates and then the *storeToCassandra* operation attached to the stream RDD started producing a reasonable number of writes (around 100/sec).

```
[root@ip-172-31-0-62 ~]# hadoop fs -cat /cloudcapstone/queries/G3Q2-results/G3Q2/part*
[At period: 240] Keys (day|destY): 21064, Days parsed: 78, Days' range: [2007-12-30, 2008-03-16]
[root@ip-172-31-0-62 ~]#
[root@ip-172-31-0-62 ~]# hadoop fs -cat /cloudcapstone/queries/G3Q2-results/G3Q2/part*
[At period: 360] Keys (day|destY): 31708, Days parsed: 117, Days' range: [2007-12-30, 2008-04-24]
[root@ip-172-31-0-62 ~]# du -hs /var/lib/cassandra/*
2.9G /var/lib/cassandra/commitlog
1.9G /var/lib/cassandra/data
4.0K /var/lib/cassandra/saved_caches
```

Ordering problem

For all the queries that I solved, I decided to save the result into both HDFS and Cassandra. Saving to HDFS, it was easier to produce results which have a custom order. Instead, while streaming updates into Cassandra with the connector, I was forced to use the natural primary/cluster key for the destination tables, and this prevents defining the performance indicator as a key to allow ordering data in the queries (if you don't know how to obtain ordered results into Cassandra, please see my Task 1 report).

Later on, I found a solution to this problem. At savepoints, I used to store the data in a table *_partial*, then having one last savepoint, right after the job is done, to store the final results into the final table, which has a different primary key and allows ordering. It worked well for the ranked queries, G2Q1 (see the two tables' definition and the results in the screenshot) and G2Q2.

```
cqlsh:cloudcapstone> select * from spark_g2q1_airport_carriers_depart_perf_partial WHERE airport = 'SRQ' ORDER BY ontimedepartureperf desc limit 10;
InvalidRequest: code=2200 [Invalid query] message="Order by is currently only supported on the clustered columns of the PRIMARY KEY, got ontimedepartureperf"
cqlsh:cloudcapstone> select * from spark_g2q1_airport_carriers_depart_perf WHERE airport = 'SRQ' ORDER BY ontimedepartureperf desc limit 10;
+-----+-----+-----+-----+-----+-----+-----+-----+-----+-----+
| airport | ontimedepartureperf | airline |
+-----+-----+-----+-----+-----+-----+-----+-----+-----+-----+
| SRQ | 0.848777 | 19707 |
| SRQ | 0.835855 | 20312 |
| SRQ | 0.803597 | 20374 |
| SRQ | 0.777255 | 20378 |
| SRQ | 0.77466 | 20211 |
| SRQ | 0.75844 | 19386 |
| SRQ | 0.753436 | 19805 |
| SRQ | 0.744 | 20295 |
| SRQ | 0.732191 | 20437 |
| SRQ | 0.723404 | 20398 |
```

```
1 -- final table, clustering ordering
2 create table spark_g2q1_airport_carriers_depart_perf (
3   airport text, city text, airline text, carrier text, countofflights bigint,
4   ontimedepartureperf float, avgdeparturedelayinminutes float,
5   primary key ((airport), ontimedepartureperf, airline))
6   with clustering order by (ontimedepartureperf desc);
7
8 -- partial table
9 create table spark_g2q1_airport_carriers_depart_perf_partial (
10  airport text, city text, airline text, carrier text, countofflights bigint,
11  ontimedepartureperf float, avgdeparturedelayinminutes float,
12  primary key ((airport), airline));
```

Final Considerations

Comparing my batch solution (Task 1) against my streaming pipeline solution (Task 2), one main difference: the pipeline gives early results, even if partial, and produces a final result which is exactly as the batch one, in almost the same processing time.

We can incur in such a situation, when the pipeline final results are not matching the batch ones, i.e. if we missed some messages or interrupted the stream earlier. But the batch still suffers from the pain of waiting till the end to see a result that, maybe, is not significantly different from a partial one produced by the pipeline. Mostly it depends on what is our goal.

For the Tom’s problem, both the batch and the pipeline give the best solution at the end of the process. And I could say my Spark streaming solution was agile enough to process the data set with the same performance or even better than my batch solution. Comparing the time, on the same cluster, I’ve got 1.5h for the batch and 2h for the pipeline but I cannot state which one is faster, because it depends on the efficiency of the solution.

My final opinion in points: (just comparing the two solutions, not Spark vs. Hadoop ;)

- Hadoop MR jobs (via Pig or other tools) over YARN/Tez are easier to design and to code in practice ... **+1 batch**
- Stateful, all-in-memory computations are complex but feasible (thanks to Spark), if well designed ... **+1 stream**
- The pipeline gives (partial) results faster ... **+1 stream**
- I had to put much more effort to setup and maintain the pipeline than maintaining the batch platform ... **+1 batch**

No loser, both win.

Demo video

A video showing my work for the Task 2 is on YouTube, https://youtu.be/0bO5Y_G_23w

If you are interested, there is also an extra video of exploring the source code of my solution, <https://youtu.be/Xp7kHObxGV4>

For any problem, or further question, please email to luiginoto@gmail.com.

Appendix - Full results

G3Q2

```
cqlsh:cloudcapstone> -- TASK 2 Queries
cqlsh:cloudcapstone> --BOS → ATL → LAX, 03/04/2008
cqlsh:cloudcapstone> select date1, deptime1, date2, deptime2, origx, desty, destz, arrdelay1, arrdelay2, origcityx, destcityy, destcityz, flightnum1, flightnum2 from cloudcapston
HERE origx = 'BOS' and desty='ATL' and destz = 'LAX' AND date1 = '2008-04-03' order BY origx, desty, destz LIMIT 1000 ALLOW FILTERING;
```

| date1 | deptime1 | date2 | deptime2 | origx | desty | destz | arrdelay1 | arrdelay2 | origcityx | destcityy | destcityz | flightnum1 | flightnum2 |
|------------|----------|------------|----------|-------|-------|-------|-----------|-----------|------------|-------------|-----------------|------------|------------|
| 2008-04-03 | 0600 | 2008-04-05 | 1852 | BOS | ATL | LAX | 7 | -2 | Boston, MA | Atlanta, GA | Los Angeles, CA | 270 | 40 |

```
(1 rows)
cqlsh:cloudcapstone>
cqlsh:cloudcapstone> --PHX → JFK → MSP, 07/09/2008
cqlsh:cloudcapstone> select date1, deptime1, date2, deptime2, origx, desty, destz, arrdelay1, arrdelay2, origcityx, destcityy, destcityz, flightnum1, flightnum2 from cloudcapston
HERE origx = 'PHX' and desty='JFK' and destz = 'MSP' AND date1 = '2008-09-07' order BY origx, desty, destz LIMIT 1000 ALLOW FILTERING;
```

| date1 | deptime1 | date2 | deptime2 | origx | desty | destz | arrdelay1 | arrdelay2 | origcityx | destcityy | destcityz | flightnum1 | flightnum2 |
|------------|----------|------------|----------|-------|-------|-------|-----------|-----------|-------------|--------------|--------------------------|------------|------------|
| 2008-09-07 | 1130 | 2008-09-09 | 1750 | PHX | JFK | MSP | -25 | -17 | Phoenix, AZ | New York, NY | Minneapolis/St. Paul, MN | 178 | 609 |

```
(1 rows)
cqlsh:cloudcapstone>
cqlsh:cloudcapstone> --DFW → STL → ORD, 24/01/2008
cqlsh:cloudcapstone> select date1, deptime1, date2, deptime2, origx, desty, destz, arrdelay1, arrdelay2, origcityx, destcityy, destcityz, flightnum1, flightnum2 from cloudcapston
HERE origx = 'DFW' and desty='STL' and destz = 'ORD' AND date1 = '2008-01-24' order BY origx, desty, destz LIMIT 1000 ALLOW FILTERING;
```

| date1 | deptime1 | date2 | deptime2 | origx | desty | destz | arrdelay1 | arrdelay2 | origcityx | destcityy | destcityz | flightnum1 | flightnum2 |
|------------|----------|------------|----------|-------|-------|-------|-----------|-----------|---------------------|---------------|-------------|------------|------------|
| 2008-01-24 | 0705 | 2008-01-26 | 1655 | DFW | STL | ORD | -14 | -5 | Dallas/Ft.Worth, TX | St. Louis, MO | Chicago, IL | 1336 | 2245 |

```
(1 rows)
cqlsh:cloudcapstone>
cqlsh:cloudcapstone> --LAX → MIA → LAX, 16/05/2008
cqlsh:cloudcapstone> select date1, deptime1, date2, deptime2, origx, desty, destz, arrdelay1, arrdelay2, origcityx, destcityy, destcityz, flightnum1, flightnum2 from cloudcapston
HERE origx = 'LAX' and desty='MIA' and destz = 'LAX' AND date1 = '2008-05-16' order BY origx, desty, destz LIMIT 1000 ALLOW FILTERING;
```

| date1 | deptime1 | date2 | deptime2 | origx | desty | destz | arrdelay1 | arrdelay2 | origcityx | destcityy | destcityz | flightnum1 | flightnum2 |
|------------|----------|------------|----------|-------|-------|-------|-----------|-----------|-----------------|-----------|-----------------|------------|------------|
| 2008-05-16 | 0820 | 2008-05-18 | 1930 | LAX | MIA | LAX | 10 | -19 | Los Angeles, CA | Miami, FL | Los Angeles, CA | 280 | 456 |

G2Q1

```
cqlsh:cloudcapstone> -- SRQ
cqlsh:cloudcapstone> select * from spark_g2q1_airport_carriers_depart_perf WHERE airport = 'SRQ' ORDER BY ontimedepartureperf desc limit 10;
```

| airport | ontimedepartureperf | airline | avgdeparturedelayinminutes | carrier | city | countofflights |
|---------|---------------------|---------|----------------------------|---------|------------------------|----------------|
| SRQ | 0.848777 | 19707 | 6.15765 | EA | Sarasota/Bradenton, FL | 5601 |
| SRQ | 0.835855 | 20312 | 4.83888 | TZ | Sarasota/Bradenton, FL | 1322 |
| SRQ | 0.803597 | 20374 | 5.12931 | XE | Sarasota/Bradenton, FL | 2057 |
| SRQ | 0.777355 | 20378 | 6.67745 | YV | Sarasota/Bradenton, FL | 1051 |
| SRQ | 0.77466 | 20211 | 5.6679 | TW | Sarasota/Bradenton, FL | 12936 |
| SRQ | 0.75844 | 19386 | 6.22155 | NW | Sarasota/Bradenton, FL | 10693 |
| SRQ | 0.753436 | 19805 | 4.88136 | AA | Sarasota/Bradenton, FL | 10042 |
| SRQ | 0.744 | 20295 | 7.288 | ML (1) | Sarasota/Bradenton, FL | 625 |
| SRQ | 0.732191 | 20437 | 9.01343 | FL | Sarasota/Bradenton, FL | 8114 |
| SRQ | 0.723404 | 20398 | 7.66667 | MQ | Sarasota/Bradenton, FL | 423 |

(10 rows)

```
cqlsh:cloudcapstone> -- CMH
cqlsh:cloudcapstone> select * from spark_g2q1_airport_carriers_depart_perf WHERE airport = 'CMH' ORDER BY ontimedepartureperf desc limit 10;
```

| airport | ontimedepartureperf | airline | avgdeparturedelayinminutes | carrier | city | countofflights |
|---------|---------------------|---------|----------------------------|---------|--------------|----------------|
| CMH | 0.855478 | 20295 | 4.26729 | ML (1) | Columbus, OH | 1287 |
| CMH | 0.855457 | 19707 | 5.23118 | EA | Columbus, OH | 4490 |
| CMH | 0.798942 | 20404 | 7.05843 | DH | Columbus, OH | 4946 |
| CMH | 0.765077 | 19805 | 4.53355 | AA | Columbus, OH | 40890 |
| CMH | 0.762421 | 20378 | 9.5682 | YV | Columbus, OH | 1107 |
| CMH | 0.76186 | 19386 | 5.40979 | NW | Columbus, OH | 58479 |
| CMH | 0.704846 | 20366 | 13.86344 | EV | Columbus, OH | 227 |
| CMH | 0.70429 | 20211 | 6.42525 | TW | Columbus, OH | 34801 |
| CMH | 0.696252 | 20374 | 11.81821 | XE | Columbus, OH | 14275 |
| CMH | 0.688267 | 20417 | 12.65506 | OH | Columbus, OH | 13040 |

(10 rows)

```
cqlsh:cloudcapstone> -- JFK
cqlsh:cloudcapstone> select * from spark_g2q1_airport_carriers_depart_perf WHERE airport = 'JFK' ORDER BY ontimedepartureperf desc limit 10;
```

| airport | ontimedepartureperf | airline | avgdeparturedelayinminutes | carrier | city | countofflights |
|---------|---------------------|---------|----------------------------|---------|--------------|----------------|
| JFK | 0.711221 | 20374 | 10.60506 | XE | New York, NY | 1818 |
| JFK | 0.697051 | 19707 | 13.24746 | EA | New York, NY | 6001 |
| JFK | 0.693871 | 20404 | 11.26115 | DH | New York, NY | 12449 |
| JFK | 0.635885 | 19977 | 6.30787 | UA | New York, NY | 94341 |
| JFK | 0.633045 | 19386 | 13.11643 | NW | New York, NY | 16765 |
| JFK | 0.620712 | 20384 | 10.99403 | PA (1) | New York, NY | 43724 |
| JFK | 0.613699 | 20417 | 17.71196 | OH | New York, NY | 70326 |
| JFK | 0.60773 | 20378 | 18.67873 | YV | New York, NY | 4321 |
| JFK | 0.580879 | 20398 | 14.6947 | MQ | New York, NY | 58699 |
| JFK | 0.580197 | 19704 | 8.67044 | CO | New York, NY | 4676 |

(10 rows)

```
cqlsh:cloudcapstone> -- SEA
cqlsh:cloudcapstone> select * from spark_g2q1_airport_carriers_depart_perf WHERE airport = 'SEA' ORDER BY ontimedepartureperf desc limit 10;
```

| airport | ontimedepartureperf | airline | avgdeparturedelayinminutes | carrier | city | countofflights |
|---------|---------------------|---------|----------------------------|---------|-------------|----------------|
| SEA | 0.832117 | 20378 | 7.30657 | YV | Seattle, WA | 548 |
| SEA | 0.748656 | 20304 | 4.70382 | OO | Seattle, WA | 27341 |
| SEA | 0.689221 | 20312 | 8.38985 | TZ | Seattle, WA | 2542 |
| SEA | 0.680443 | 19707 | 11.46716 | EA | Seattle, WA | 4065 |
| SEA | 0.666212 | 19386 | 7.21636 | NW | Seattle, WA | 118000 |
| SEA | 0.651852 | 20366 | 8.37037 | EV | Seattle, WA | 270 |
| SEA | 0.651282 | 20404 | 12.14359 | DH | Seattle, WA | 195 |
| SEA | 0.647513 | 19690 | 8.16441 | HA | Seattle, WA | 3759 |
| SEA | 0.630984 | 19805 | 7.32507 | AA | Seattle, WA | 123000 |
| SEA | 0.628542 | 19991 | 9.17828 | HP | Seattle, WA | 46051 |

(10 rows)

```
cqlsh:cloudcapstone> -- BOS
cqlsh:cloudcapstone> select * from spark_g2q1_airport_carriers_depart_perf WHERE airport = 'BOS' ORDER BY ontimedepartureperf desc limit 10;
```

| airport | ontimedepartureperf | airline | avgdeparturedelayinminutes | carrier | city | countofflights |
|---------|---------------------|---------|----------------------------|---------|------------|----------------|
| BOS | 0.835209 | 20295 | 5.49357 | ML (1) | Boston, MA | 1244 |
| BOS | 0.798593 | 20384 | 4.53687 | PA (1) | Boston, MA | 25168 |
| BOS | 0.778326 | 19707 | 7.99648 | EA | Boston, MA | 28393 |
| BOS | 0.772821 | 20312 | 6.25643 | TZ | Boston, MA | 3385 |
| BOS | 0.718062 | 20374 | 9.7467 | XE | Boston, MA | 2270 |
| BOS | 0.714126 | 20437 | 13.16743 | FL | Boston, MA | 41959 |
| BOS | 0.711592 | 20363 | 13.90373 | 9E | Boston, MA | 2493 |
| BOS | 0.687432 | 20417 | 14.18676 | OH | Boston, MA | 53758 |
| BOS | 0.678562 | 19386 | 8.23316 | NW | Boston, MA | 142158 |
| BOS | 0.671502 | 20404 | 12.34528 | DH | Boston, MA | 18889 |

(10 rows)

G2Q2

```
cqlsh:cloudcapstone> -- SRQ
cqlsh:cloudcapstone> select * from spark_g2q2_airport_routes_depart_perf WHERE origin = 'SRQ' ORDER BY ontimedepartureperf desc limit 10;
```

| origin | ontimedepartureperf | dest | avgdeparturedelayinminutes | countofflights | destcity | origcity |
|--------|---------------------|------|----------------------------|----------------|--------------------------|------------------------|
| SRQ | 1 | EYW | 0 | 1 | Key West, FL | Sarasota/Bradenton, FL |
| SRQ | 0.846154 | MSP | 6.09615 | 184 | Minneapolis/St. Paul, MN | Sarasota/Bradenton, FL |
| SRQ | 0.835326 | TPA | 2.18126 | 2593 | Tampa, FL | Sarasota/Bradenton, FL |
| SRQ | 0.822967 | DCA | 4.7177 | 209 | Washington, DC | Sarasota/Bradenton, FL |
| SRQ | 0.820359 | BNA | 2.81437 | 334 | Nashville, TN | Sarasota/Bradenton, FL |
| SRQ | 0.817797 | MEM | 1.86758 | 944 | Memphis, TN | Sarasota/Bradenton, FL |
| SRQ | 0.780379 | IAH | 3.52806 | 2905 | Houston, TX | Sarasota/Bradenton, FL |
| SRQ | 0.775709 | RSW | 4.6302 | 12515 | Ft. Myers, FL | Sarasota/Bradenton, FL |
| SRQ | 0.773573 | STL | 5.91417 | 4730 | St. Louis, MO | Sarasota/Bradenton, FL |
| SRQ | 0.760548 | BWI | 6.05696 | 1896 | Baltimore, MD | Sarasota/Bradenton, FL |

(10 rows)

```
cqlsh:cloudcapstone> -- CMH
cqlsh:cloudcapstone> select * from spark_g2q2_airport_routes_depart_perf WHERE origin = 'CMH' ORDER BY ontimedepartureperf desc limit 10;
```

| origin | ontimedepartureperf | dest | avgdeparturedelayinminutes | countofflights | destcity | origcity |
|--------|---------------------|------|----------------------------|----------------|----------------------|--------------|
| CMH | 1 | AUS | 0 | 3 | Austin, TX | Columbus, OH |
| CMH | 1 | OMA | -5 | 1 | Omaha, NE | Columbus, OH |
| CMH | 0.84874 | SDF | 2.7395 | 119 | Louisville, KY | Columbus, OH |
| CMH | 0.814375 | CLE | 3.88003 | 7235 | Cleveland, OH | Columbus, OH |
| CMH | 0.768627 | SLC | 6.92353 | 510 | Salt Lake City, UT | Columbus, OH |
| CMH | 0.762457 | DTW | 5.82289 | 34137 | Detroit, MI | Columbus, OH |
| CMH | 0.757574 | DFW | 5.53413 | 27761 | Dallas/Ft. Worth, TX | Columbus, OH |
| CMH | 0.753537 | MEM | 5.92517 | 5372 | Memphis, TN | Columbus, OH |
| CMH | 0.752987 | BNA | 4.64757 | 9048 | Nashville, TN | Columbus, OH |
| CMH | 0.75285 | IAD | 7.38501 | 5790 | Washington, DC | Columbus, OH |

(10 rows)

```
cqlsh:cloudcapstone> -- JFK
cqlsh:cloudcapstone> select * from spark_g2q2_airport_routes_depart_perf WHERE origin = 'JFK' ORDER BY ontimedepartureperf desc limit 10;
```

| origin | ontimedepartureperf | dest | avgdeparturedelayinminutes | countofflights | destcity | origcity |
|--------|---------------------|------|----------------------------|----------------|---------------------------|--------------|
| JFK | 1 | ABQ | 0 | 1 | Albuquerque, NM | New York, NY |
| JFK | 1 | ANC | 0 | 27 | Anchorage, AK | New York, NY |
| JFK | 1 | ISP | 0 | 1 | Islip, NY | New York, NY |
| JFK | 1 | MYR | 0 | 1 | Myrtle Beach, SC | New York, NY |
| JFK | 1 | SWF | -0.5 | 2 | Newburgh/Poughkeepsie, NY | New York, NY |
| JFK | 0.830508 | BGR | 6.55085 | 118 | Bangor, ME | New York, NY |
| JFK | 0.828025 | DAB | 7.56688 | 157 | Daytona Beach, FL | New York, NY |
| JFK | 0.821705 | CHS | 7.38953 | 1032 | Charleston, SC | New York, NY |
| JFK | 0.814751 | PMS | 7.90566 | 583 | Pensacola, FL | New York, NY |
| JFK | 0.804264 | SAV | 8.08915 | 1032 | Savannah, GA | New York, NY |

(10 rows)

```
cqlsh:cloudcapstone> -- SEA
cqlsh:cloudcapstone> select * from spark_g2q2_airport_routes_depart_perf WHERE origin = 'SEA' ORDER BY ontimedepartureperf desc limit 10;
```

| origin | ontimedepartureperf | dest | avgdeparturedelayinminutes | countofflights | destcity | origcity |
|--------|---------------------|------|----------------------------|----------------|------------------------------|-------------|
| SEA | 1 | EUG | 0 | 1 | Eugene, OR | Seattle, WA |
| SEA | 0.802181 | PSC | 4.09717 | 2017 | Pasco/Kennewick/Richland, WA | Seattle, WA |
| SEA | 0.764151 | ICT | 11.90566 | 106 | Wichita, KS | Seattle, WA |
| SEA | 0.695017 | MEM | 4.96892 | 11358 | Memphis, TN | Seattle, WA |
| SEA | 0.68033 | PDX | 7.48523 | 74114 | Portland, OR | Seattle, WA |
| SEA | 0.678473 | CLE | 6.63239 | 2908 | Cleveland, OH | Seattle, WA |
| SEA | 0.670561 | DTW | 6.40699 | 26278 | Detroit, MI | Seattle, WA |
| SEA | 0.667544 | SNA | 6.16402 | 38745 | Santa Ana, CA | Seattle, WA |
| SEA | 0.659743 | MSP | 7.38041 | 44628 | Minneapolis/St. Paul, MN | Seattle, WA |
| SEA | 0.652543 | GEG | 7.90643 | 79066 | Spokane, WA | Seattle, WA |

(10 rows)

```
cqlsh:cloudcapstone> -- BOS
cqlsh:cloudcapstone> select * from spark_g2q2_airport_routes_depart_perf WHERE origin = 'BOS' ORDER BY ontimedepartureperf desc limit 10;
```

| origin | ontimedepartureperf | dest | avgdeparturedelayinminutes | countofflights | destcity | origcity |
|--------|---------------------|------|----------------------------|----------------|-------------------------------|------------|
| BOS | 1 | ONT | 0 | 1 | Ontario/San Bernardino, CA | Boston, MA |
| BOS | 1 | SWF | 0 | 1 | Newburgh/Poughkeepsie, NY | Boston, MA |
| BOS | 0.813449 | AUS | 3.57122 | 1383 | Austin, TX | Boston, MA |
| BOS | 0.787836 | LGA | 3.95273 | 163647 | New York, NY | Boston, MA |
| BOS | 0.78534 | MSY | 6.13613 | 573 | New Orleans, LA | Boston, MA |
| BOS | 0.770202 | MYR | 8.61111 | 396 | Myrtle Beach, SC | Boston, MA |
| BOS | 0.754503 | MDW | 8.05957 | 7051 | Chicago, IL | Boston, MA |
| BOS | 0.712457 | MKE | 6.88105 | 3203 | Milwaukee, WI | Boston, MA |
| BOS | 0.704944 | SAV | 11.78947 | 627 | Savannah, GA | Boston, MA |
| BOS | 0.703515 | PHF | 14.49925 | 3329 | Newport News/Williamsburg, VA | Boston, MA |

(10 rows)

G2Q2 (from HDFS)

```
[root@ip-172-31-0-62 ~]# hadoop fs -cat /cloudcapstone/spark-results/G2Q2/part* | grep "SRQ-" | head
(SRQ-EYW,1.0,1,0.0)
(SRQ-MSP,0.84615386,104,6.0961537)
(SRQ-TPA,0.8353259,2593,2.181257)
(SRQ-DCA,0.8229665,209,4.717703)
(SRQ-BNA,0.8203593,334,2.8143709)
(SRQ-MEM,0.8177966,944,1.8675847)
(SRQ-IAH,0.78037876,2905,3.5280552)
(SRQ-RSW,0.77570915,12515,4.6302023)
(SRQ-STL,0.773573,4730,5.914165)
(SRQ-BWI,0.7605485,1896,6.056961)
[root@ip-172-31-0-62 ~]# hadoop fs -cat /cloudcapstone/spark-results/G2Q2/part* | grep "CMH-" | head
(CMH-AUS,1.0,3,0.0)
(CMH-OMA,1.0,1,-5.0)
(CMH-SDF,0.8487395,119,2.7394958)
(CMH-CLE,0.8143746,7235,3.8800275)
(CMH-SLC,0.76862746,510,6.9235296)
(CMH-DTW,0.76245713,34137,5.8228908)
(CMH-DFW,0.75757354,27761,5.5341296)
(CMH-MEM,0.7535369,5372,5.925168)
(CMH-BNA,0.7529867,9040,4.6475673)
(CMH-IAD,0.75284976,5790,7.3050094)
[root@ip-172-31-0-62 ~]# hadoop fs -cat /cloudcapstone/spark-results/G2Q2/part* | grep "JFK-" | head
(JFK-ABQ,1.0,1,0.0)
(JFK-ANC,1.0,27,0.0)
(JFK-ISP,1.0,1,0.0)
(JFK-MYR,1.0,1,0.0)
(JFK-SWF,1.0,2,-0.5)
(JFK-BGR,0.8305085,118,6.5508475)
(JFK-DAB,0.82802546,157,7.566879)
(JFK-CHS,0.8217054,1032,7.389535)
(JFK-PNS,0.81475127,583,7.905661)
(JFK-SAV,0.8042636,1032,8.089148)
[root@ip-172-31-0-62 ~]# hadoop fs -cat /cloudcapstone/spark-results/G2Q2/part* | grep "SEA-" | head
(SEA-EUG,1.0,1,0.0)
(SEA-PSC,0.8021815,2017,4.0971737)
(SEA-ICT,0.7641509,106,11.905661)
(SEA-MEM,0.6950168,11358,4.9689226)
(SEA-PDX,0.6803302,74114,7.485227)
(SEA-CLE,0.67847323,2908,6.6323934)
(SEA-DTW,0.6705608,26278,6.4069963)
(SEA-SNA,0.6675441,38745,6.164022)
(SEA-MSP,0.65974283,44628,7.3804126)
(SEA-GEG,0.65254337,79066,7.9064326)
[root@ip-172-31-0-62 ~]# hadoop fs -cat /cloudcapstone/spark-results/G2Q2/part* | grep "BOS-" | head
(BOS-ONT,1.0,1,0.0)
(BOS-SWF,1.0,1,0.0)
(BOS-AUS,0.813449,1383,3.571222)
(BOS-LGA,0.78783613,163647,3.9527278)
(BOS-MSY,0.7853403,573,6.1361256)
(BOS-MYR,0.77020204,396,8.611111)
(BOS-MDW,0.7545029,7051,8.059566)
(BOS-MKE,0.71245724,3203,6.881049)
(BOS-SAV,0.70494413,627,11.7894745)
(BOS-PHF,0.7035146,3329,14.499249)
```

G2Q3

```
cqlsh:cloudcapstone>
cqlsh:cloudcapstone> -- LGA → BOS
cqlsh:cloudcapstone> select * from spark_g2q4_route_mean_arrival_delay WHERE origin = 'LGA' and destination = 'BOS';
```

| origin | destination | avgarrivaldelayinminutes | countofflights |
|--------|-------------|--------------------------|----------------|
| LGA | BOS | 7.66222 | 166438 |

(1 rows)

```
cqlsh:cloudcapstone> -- BOS → LGA
cqlsh:cloudcapstone> select * from spark_g2q4_route_mean_arrival_delay WHERE origin = 'BOS' and destination = 'LGA';
```

| origin | destination | avgarrivaldelayinminutes | countofflights |
|--------|-------------|--------------------------|----------------|
| BOS | LGA | 8.49688 | 163647 |

(1 rows)

```
cqlsh:cloudcapstone> -- MSP → ATL
cqlsh:cloudcapstone> select * from spark_g2q4_route_mean_arrival_delay WHERE origin = 'MSP' and destination = 'ATL';
```

| origin | destination | avgarrivaldelayinminutes | countofflights |
|--------|-------------|--------------------------|----------------|
| MSP | ATL | 12.11548 | 78108 |

(1 rows)

```
cqlsh:cloudcapstone> -- OKC → DFW
cqlsh:cloudcapstone> select * from spark_g2q4_route_mean_arrival_delay WHERE origin = 'OKC' and destination = 'DFW';
```

| origin | destination | avgarrivaldelayinminutes | countofflights |
|--------|-------------|--------------------------|----------------|
| OKC | DFW | 8.38474 | 72186 |

(1 rows)

```
cqlsh:cloudcapstone>
```